

# Safety Helmet Detection

Using YOLOv8

Noor Ali Fadhl  
Deep Learning Project

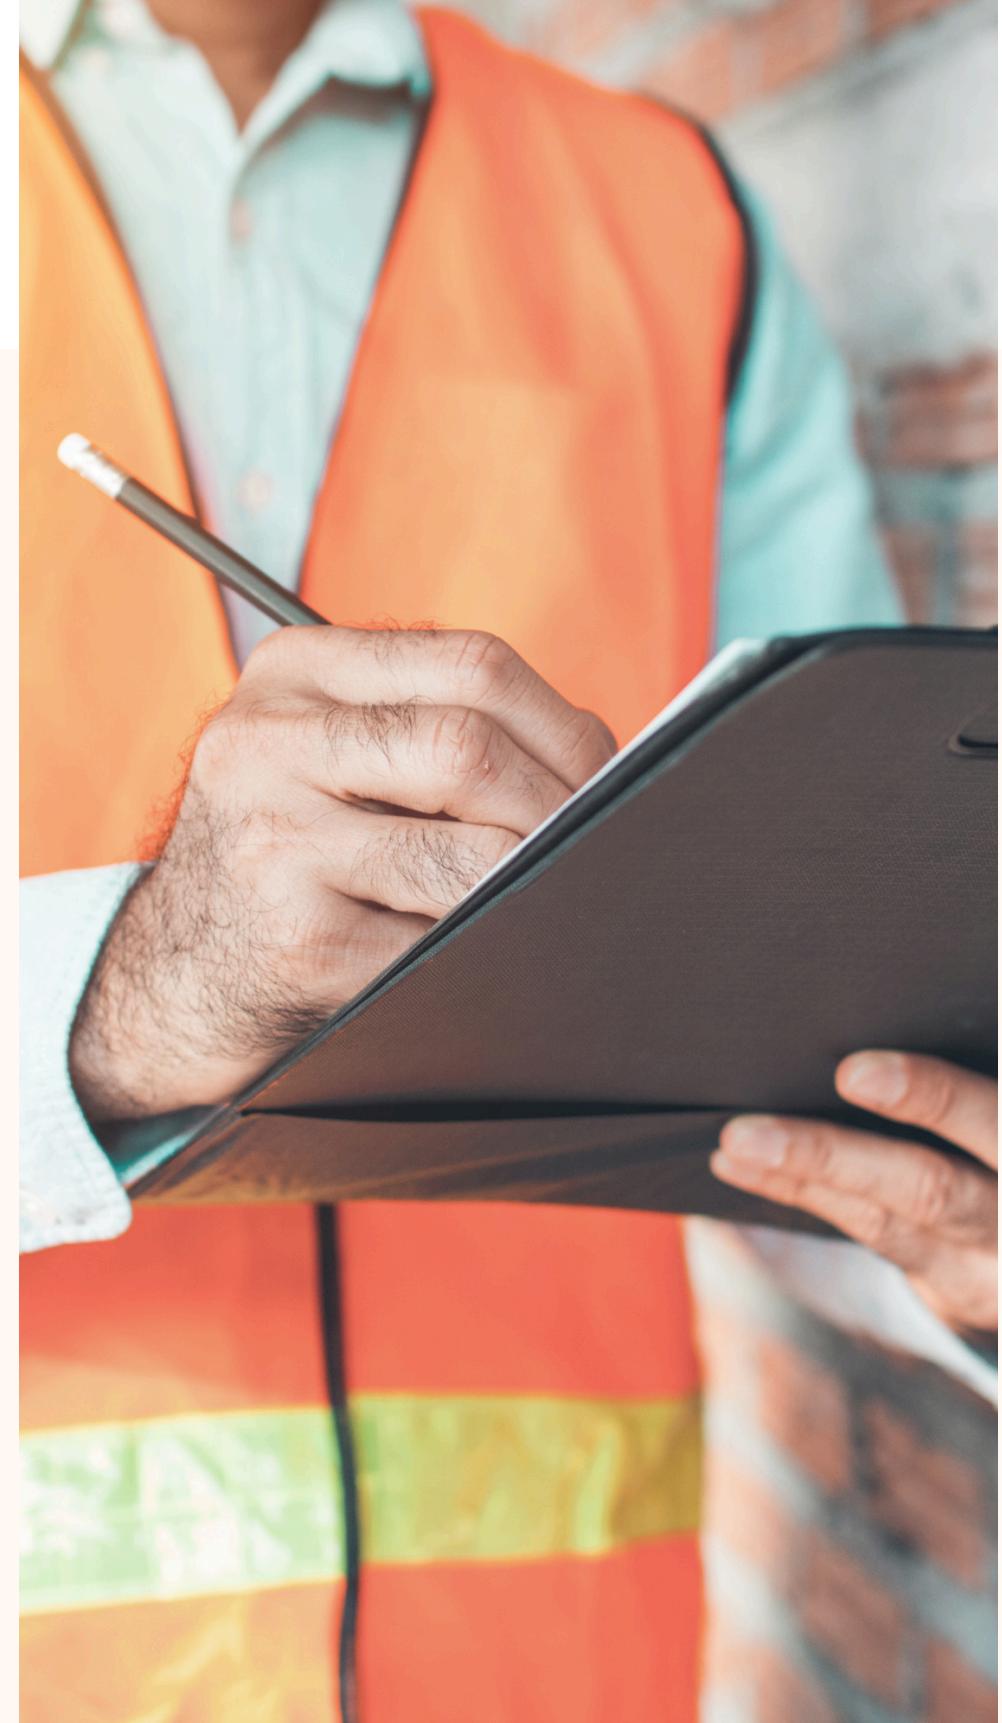


# Introduction

This project aims to improve workplace safety by using deep learning and computer vision techniques. A YOLOv8-based helmet detection system was developed to automatically detect whether workers are wearing safety helmets in real time. The system also includes a graphical user interface for easy visualization and interaction.

## Problem Statement

- Workers often fail to comply with safety helmet regulations
- Workplace accidents remain a serious safety concern
- There is a strong need for an automated and intelligent monitoring system



# Dataset

- Dataset Name: Safety Helmet Wearing Dataset
- Source: Public dataset (Kaggle)
- Total Images: ~5,000 images
- Image Type: RGB images of construction sites
- ◆ Classes
  - Helmet
  - Head (No Helmet)
  - Person
- ◆ Annotations
  - Bounding box annotations
  - YOLO format
- ◆ Data Split
  - Training set: ~70%
  - Validation set: ~30%

```
▶ import kagglehub

# Download latest version
path = kagglehub.dataset_download("andrewmvd/hard-hat-detection")

print("Path to dataset files:", path)
...
... Downloading from https://www.kaggle.com/api/v1/datasets/download/andrewmvd/hard-hat-detection?dataset\_version\_number=1...
100%|██████████| 1.22G/1.22G [00:13<00:00, 101MB/s]Extracting files...

Path to dataset files: /root/.cache/kagglehub/datasets/andrewmvd/hard-hat-detection/versions/1
```

```
import os

train_images = len(os.listdir("/content/helmet_yolo/images/train"))
val_images   = len(os.listdir("/content/helmet_yolo/images/val"))

print("Train images:", train_images)
print("Validation images:", val_images)
```

```
Train images: 4000
Validation images: 1000
```

```
▶ from collections import Counter

label_dir = "/content/helmet_yolo/labels/train"
class_counts = Counter()

classes = {0:"helmet", 1:"head", 2:"person"}

for file in os.listdir(label_dir):
    with open(os.path.join(label_dir, file)) as f:
        for line in f:
            cls = int(line.split()[0])
            class_counts[cls] += 1

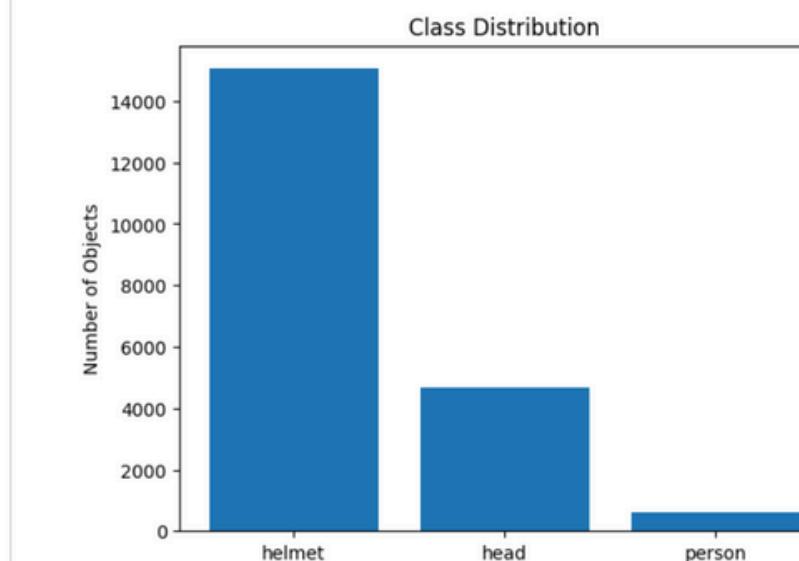
for k,v in class_counts.items():
    print(classes[k], ":", v)

...
... helmet : 15050
head : 4686
person : 597
```

```
import matplotlib.pyplot as plt

labels = [classes[k] for k in class_counts.keys()]
values = class_counts.values()

plt.bar(labels, values)
plt.title("Class Distribution")
plt.ylabel("Number of Objects")
plt.show()
```



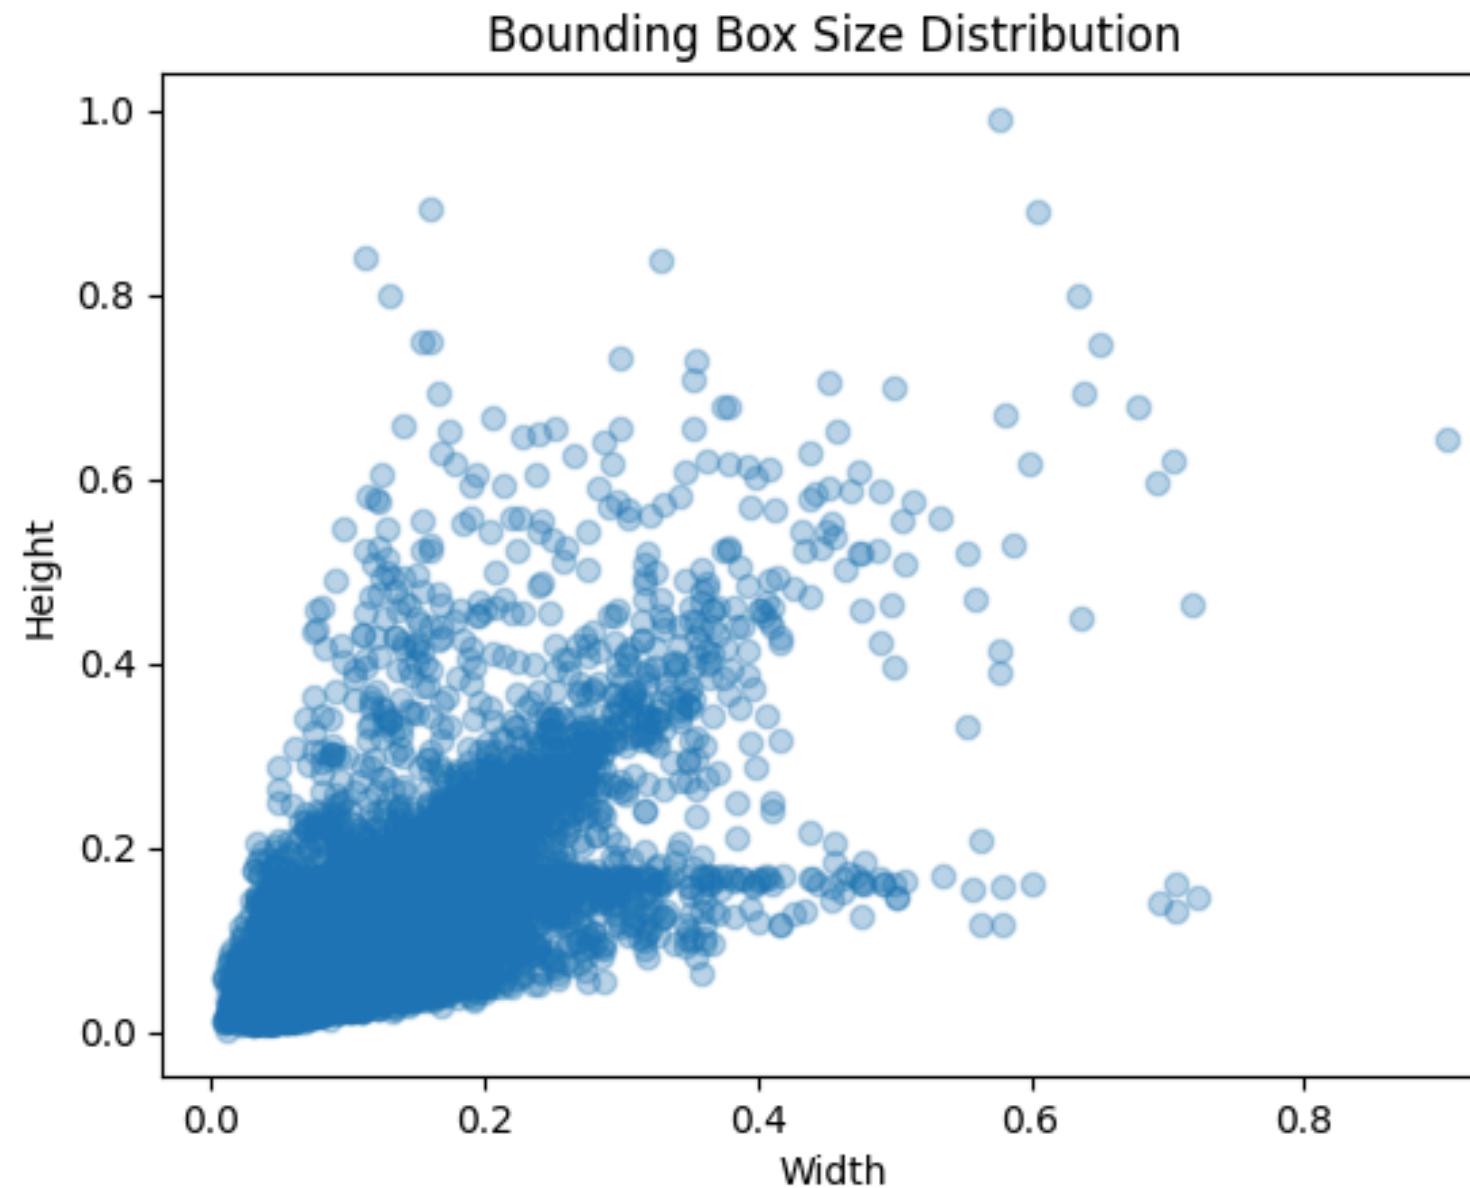
```

widths, heights = [], []

for file in os.listdir(label_dir):
    with open(os.path.join(label_dir, file)) as f:
        for line in f:
            _, x, y, w, h = map(float, line.split())
            widths.append(w)
            heights.append(h)

plt.scatter(widths, heights, alpha=0.3)
plt.xlabel("Width")
plt.ylabel("Height")
plt.title("Bounding Box Size Distribution")
plt.show()

```



```

import cv2
from ultralytics.utils.plotting import Annotator
from google.colab.patches import cv2_imshow

img_path = "/content/helmet_yolo/images/train"
lbl_path = "/content/helmet_yolo/labels/train"

sample_img = os.listdir(img_path)[0]
img = cv2.imread(os.path.join(img_path, sample_img))
h, w, _ = img.shape

annotator = Annotator(img)

with open(os.path.join(lbl_path, sample_img.replace(".png", ".txt"))) as f:
    for line in f:
        cls, x, y, bw, bh = map(float, line.split())
        x1 = int((x - bw/2) * w)
        y1 = int((y - bh/2) * h)
        x2 = int((x + bw/2) * w)
        y2 = int((y + bh/2) * h)
        annotator.box_label([x1, y1, x2, y2], {0:"helmet",1:"head",2:"person"}[int(cls)])
cv2_imshow(img)

```



# Model Architecture (YOLOv8)

- YOLOv8n (lightweight & fast)
- One-stage object detector
- Backbone + Neck + Head
- Advantages:
  - Real-time detection
  - High accuracy
  - Easy deployment

```
from ultralytics import YOLO
model = YOLO("yolov8n.pt")
model.train(
    data="/content/helmet_yolo/data.yaml",
    epochs=120,
    imgsz=640,
    batch=16,
    patience=5 # ← Early stopping
)
```

```
model = YOLO("yolov8n.pt")
model.train(
    data="/content/helmet_yolo/data.yaml",
    epochs=50,
    imgsz=640,
    batch=16,
    patience=5 # ← Early stopping
)
...
```

	50/50	3.82G	1.108	0.5183	1.041	115	640: 100%	mAP50	mAP50-95): 100%
Class		Images	Instances	Box(P	R				
all	1000	4932	0.618	0.593	0.633	0.421			

50 epochs completed in 1.403 hours.  
Optimizer stripped from /content/runs/detect/train2/weights/last.pt, 6.2MB  
Optimizer stripped from /content/runs/detect/train2/weights/best.pt, 6.2MB

Validating /content/runs/detect/train2/weights/best.pt...  
Ultralytics 8.3.241 🚀 Python-3.12.12 torch-2.9.0+cu126 CUDA:0 (Tesla T4, 15095MiB)  
Model summary (fused): 72 layers, 3,006,233 parameters, 0 gradients, 8.1 GFLOPs

Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%
all	1000	4932	0.618	0.593	0.633	0.421
helmet	925	3738	0.941	0.903	0.961	0.643
background	75	1251	0.917	0.875	0.935	0.615

# Training & Optimization

- Confidence threshold ↓ (0.2)
- IOU tuning
- Test-Time Augmentation (TTA)
  - Focus on:
    - Maximizing Recall
    - Reducing missed detections

```
import albumentations as A

# --- Paths ---
IMG_DIR = "/content/helmet_yolo/images/train"
LBL_DIR = "/content/helmet_yolo/labels/train"

OUT_IMG_DIR = "/content/helmet_yolo/images/train_aug"
OUT_LBL_DIR = "/content/helmet_yolo/labels/train_aug"

os.makedirs(OUT_IMG_DIR, exist_ok=True)
os.makedirs(OUT_LBL_DIR, exist_ok=True)

# --- Augmentation pipeline using Affine ---
transform = A.Compose([
    A.HorizontalFlip(p=0.5),
    A.VerticalFlip(p=0.3),
    A.RandomBrightnessContrast(p=0.5),
    A.Affine(translate_percent=0.05, scale=(0.9,1.1), rotate=(-15,15), p=0.7)
], bbox_params=A.BboxParams(format='yolo', label_fields=['class_labels']))

# --- Minority classes to augment ---
classes_to_augment = [1, 2] # 1: Head, 2: Person

# --- Loop through images ---
for lbl_file in os.listdir(LBL_DIR):
    lbl_path = os.path.join(LBL_DIR, lbl_file)
    img_file = lbl_file.replace(".txt", ".png") # Adjust if your images are .jpg
    img_path = os.path.join(IMG_DIR, img_file)

    if not os.path.exists(img_path):
        continue
```

```
# Read image and labels
img = cv2.imread(img_path)
with open(lbl_path) as f:
    bboxes = []
    class_labels = []
for line in f:
    cls, x, y, w, h = map(float, line.strip().split())
    if int(cls) in classes_to_augment:
        bboxes.append([x, y, w, h])
        class_labels.append(int(cls))

# Apply augmentation only if the image contains minority classes
if len(bboxes) > 0:
    for i in range(3): # Number of augmented copies per image
        augmented = transform(image=img, bboxes=bboxes, class_labels=class_labels)
        aug_img = augmented['image']
        aug_bboxes = augmented['bboxes']
        aug_labels = augmented['class_labels']

        # Save augmented image
        out_img_name = img_file.replace(".png", f"_aug{i}.png")
        cv2.imwrite(os.path.join(OUT_IMG_DIR, out_img_name), aug_img)

        # Save augmented labels
        out_lbl_name = lbl_file.replace(".txt", f"_aug{i}.txt")
        with open(os.path.join(OUT_LBL_DIR, out_lbl_name), "w") as f_out:
            for cls, bbox in zip(aug_labels, aug_bboxes):
                f_out.write(f"{cls} {' '.join(map(str, bbox))}\n")

print("✅ Dataset augmented successfully using Affine transform!")
```

```
from ultralytics import YOLO

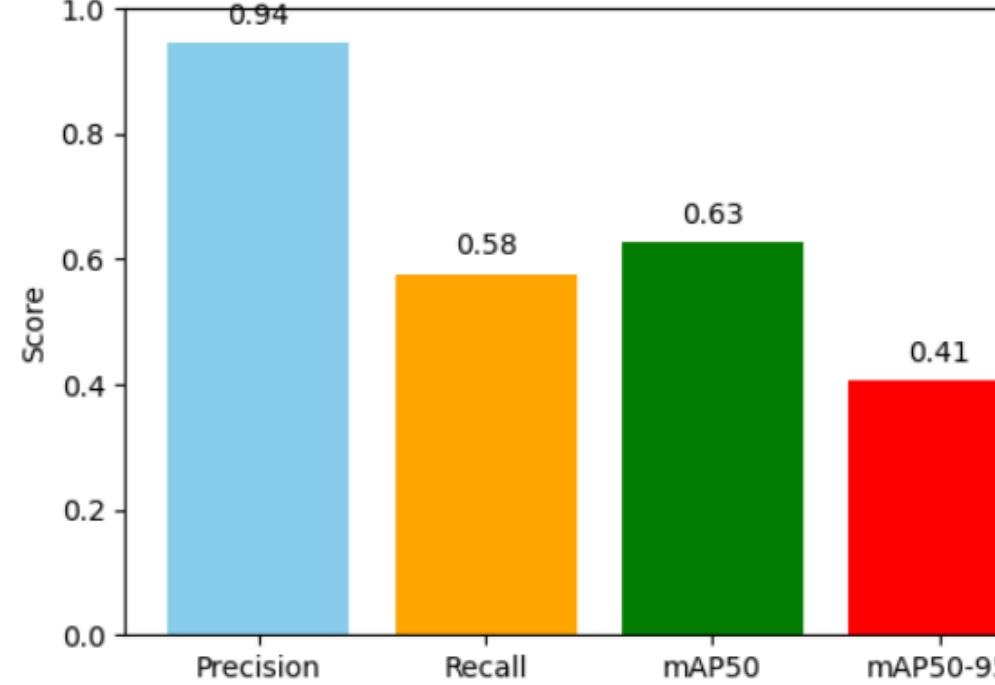
model = YOLO("/content/runs/detect/train/weights/best.pt")

model.predict(
    source="/content/helmet_yolo/images/val",
    conf=0.25,
    device=0, # GPU
    save=True
)
```

... YOLO Model Evaluation (Overall):

Metric	Value
0 Precision	0.9448
1 Recall	0.5771
2 mAP50	0.6261
3 mAP50-95	0.4067

YOLO Model Evaluation Metrics



... Ultralytics 8.3.241 Python-3.12.12 torch-2.9.0+cpu CPU (AMD EPYC 7B12)

Model summary (fused): 72 layers, 3,006,233 parameters, 0 gradients, 8.1 GFLOPS

val: Fast image access (ping: 0.0±0.0 ms, read: 35.2±14.1 MB/s, size: 266.5 KB)

val: Scanning /content/helmet\_yolo/labels/val.cache... 1000 images, 0 backgrounds, 0 corrupt: 100% 1000/1000 241.

Class	Images	Instances	Box(P	R	mAP50	mAP50-95
all	1000	5169	0.961	0.6	0.65	0.438
helmet	912	3916	0.948	0.922	0.97	0.654
head	185	1099	0.936	0.877	0.94	0.638
person	36	154	1	0	0.0389	0.0224

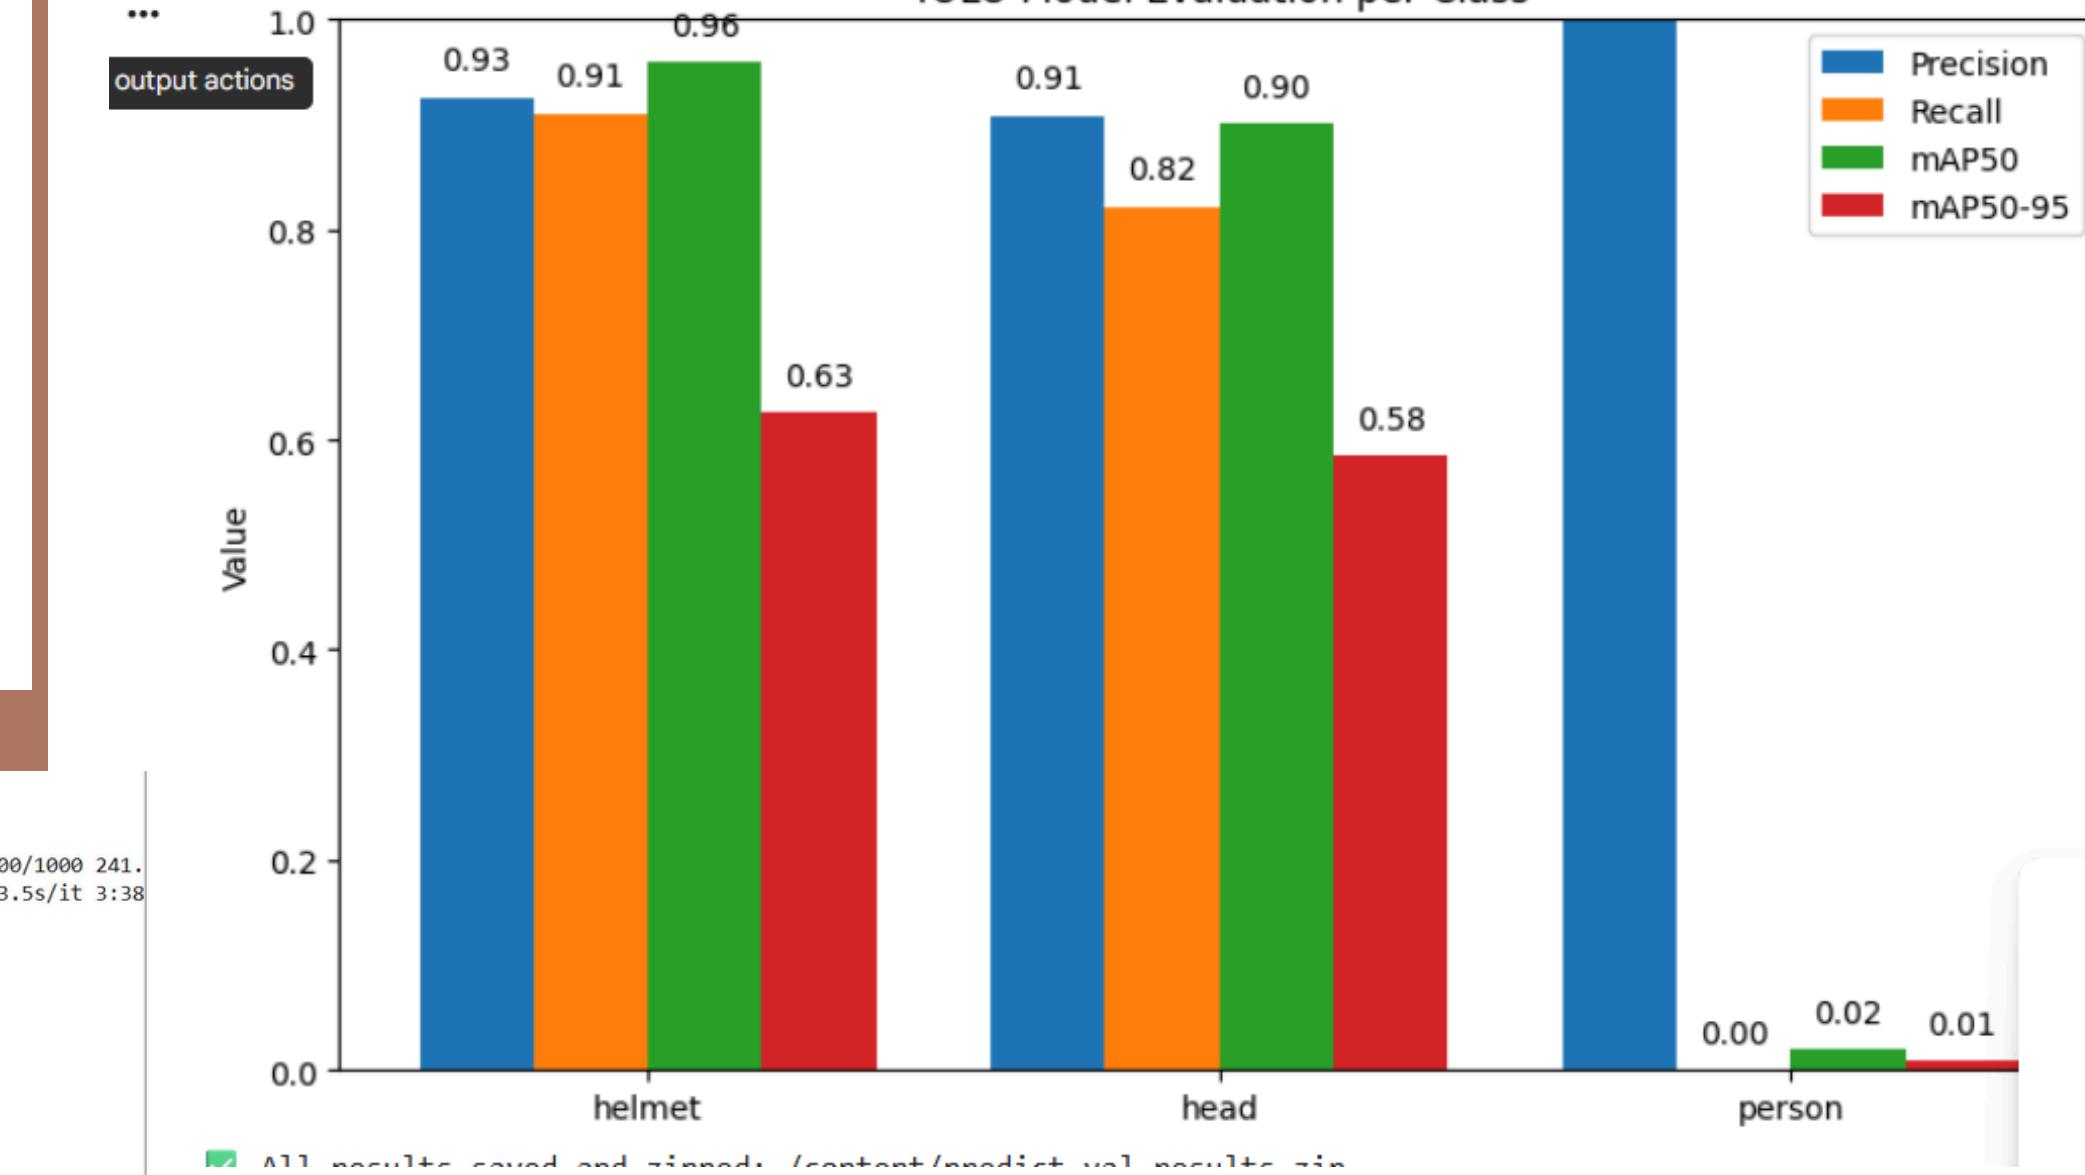
Speed: 4.3ms preprocess, 200.0ms inference, 0.0ms loss, 0.5ms postprocess per image

Results saved to /content/runs/detect/val7

Class	Precision	Recall	mAP50	mAP50-95
0 helmet	0.948017	0.922104	0.969963	0.653776
1 head	0.935891	0.876696	0.939827	0.637954
2 person	1.000000	0.000000	0.038919	0.022407

# Evaluation Metrics

YOLO Model Evaluation per Class



# GUI (User Interface)

- Built using Gradio
- Features:
  - Upload image
  - Adjust confidence & IOU
  - View predicted image
  - View evaluation table & chart
  - Download results as ZIP

```
!pip install gradio ultralytics opencv-python

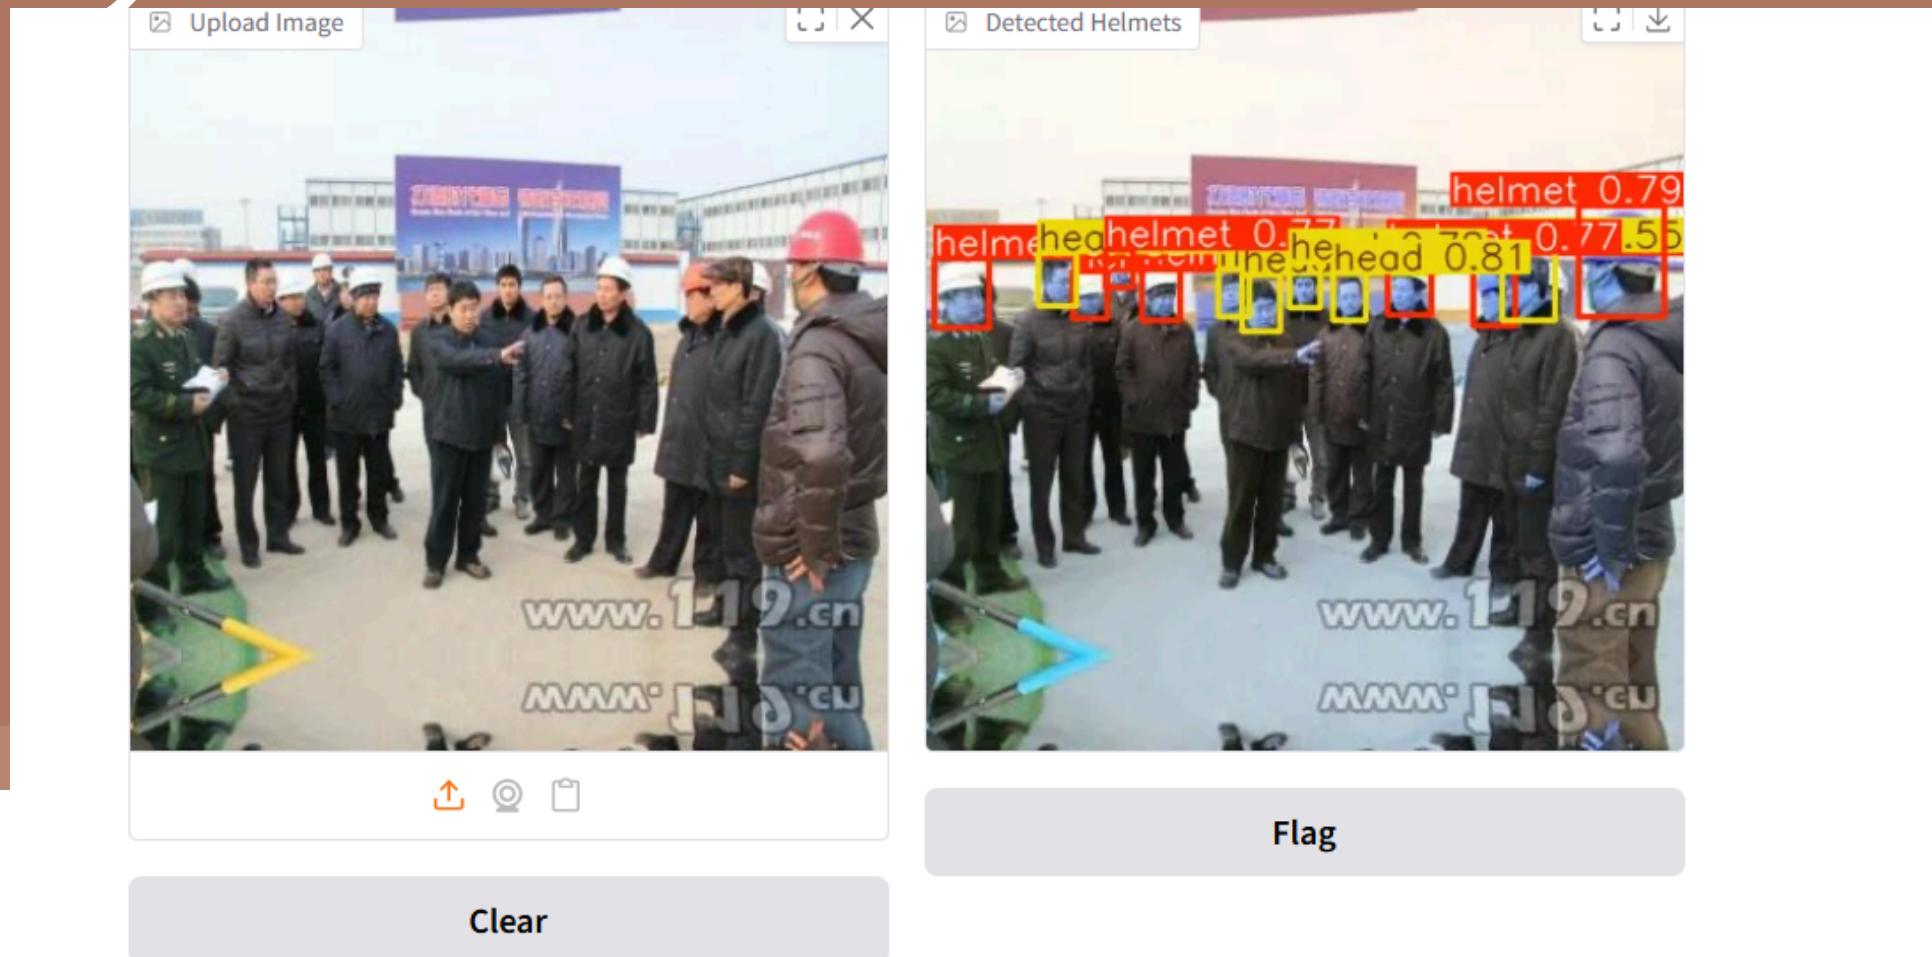
from ultralytics import YOLO
import gradio as gr

# Load your trained YOLO model
model = YOLO("/content/best (3).pt")

# Function to detect helmets on an image
def detect_helmet(image):
    results = model.predict(image, save=True) # saves image with bounding boxes
    return results[0].plot() # returns the image with detections

# Create Gradio interface
iface = gr.Interface(
    fn=detect_helmet,
    inputs=gr.Image(type="filepath", label="Upload Image"),
    outputs=gr.Image(label="Detected Helmets"),
    live=True
)

iface.launch()
```

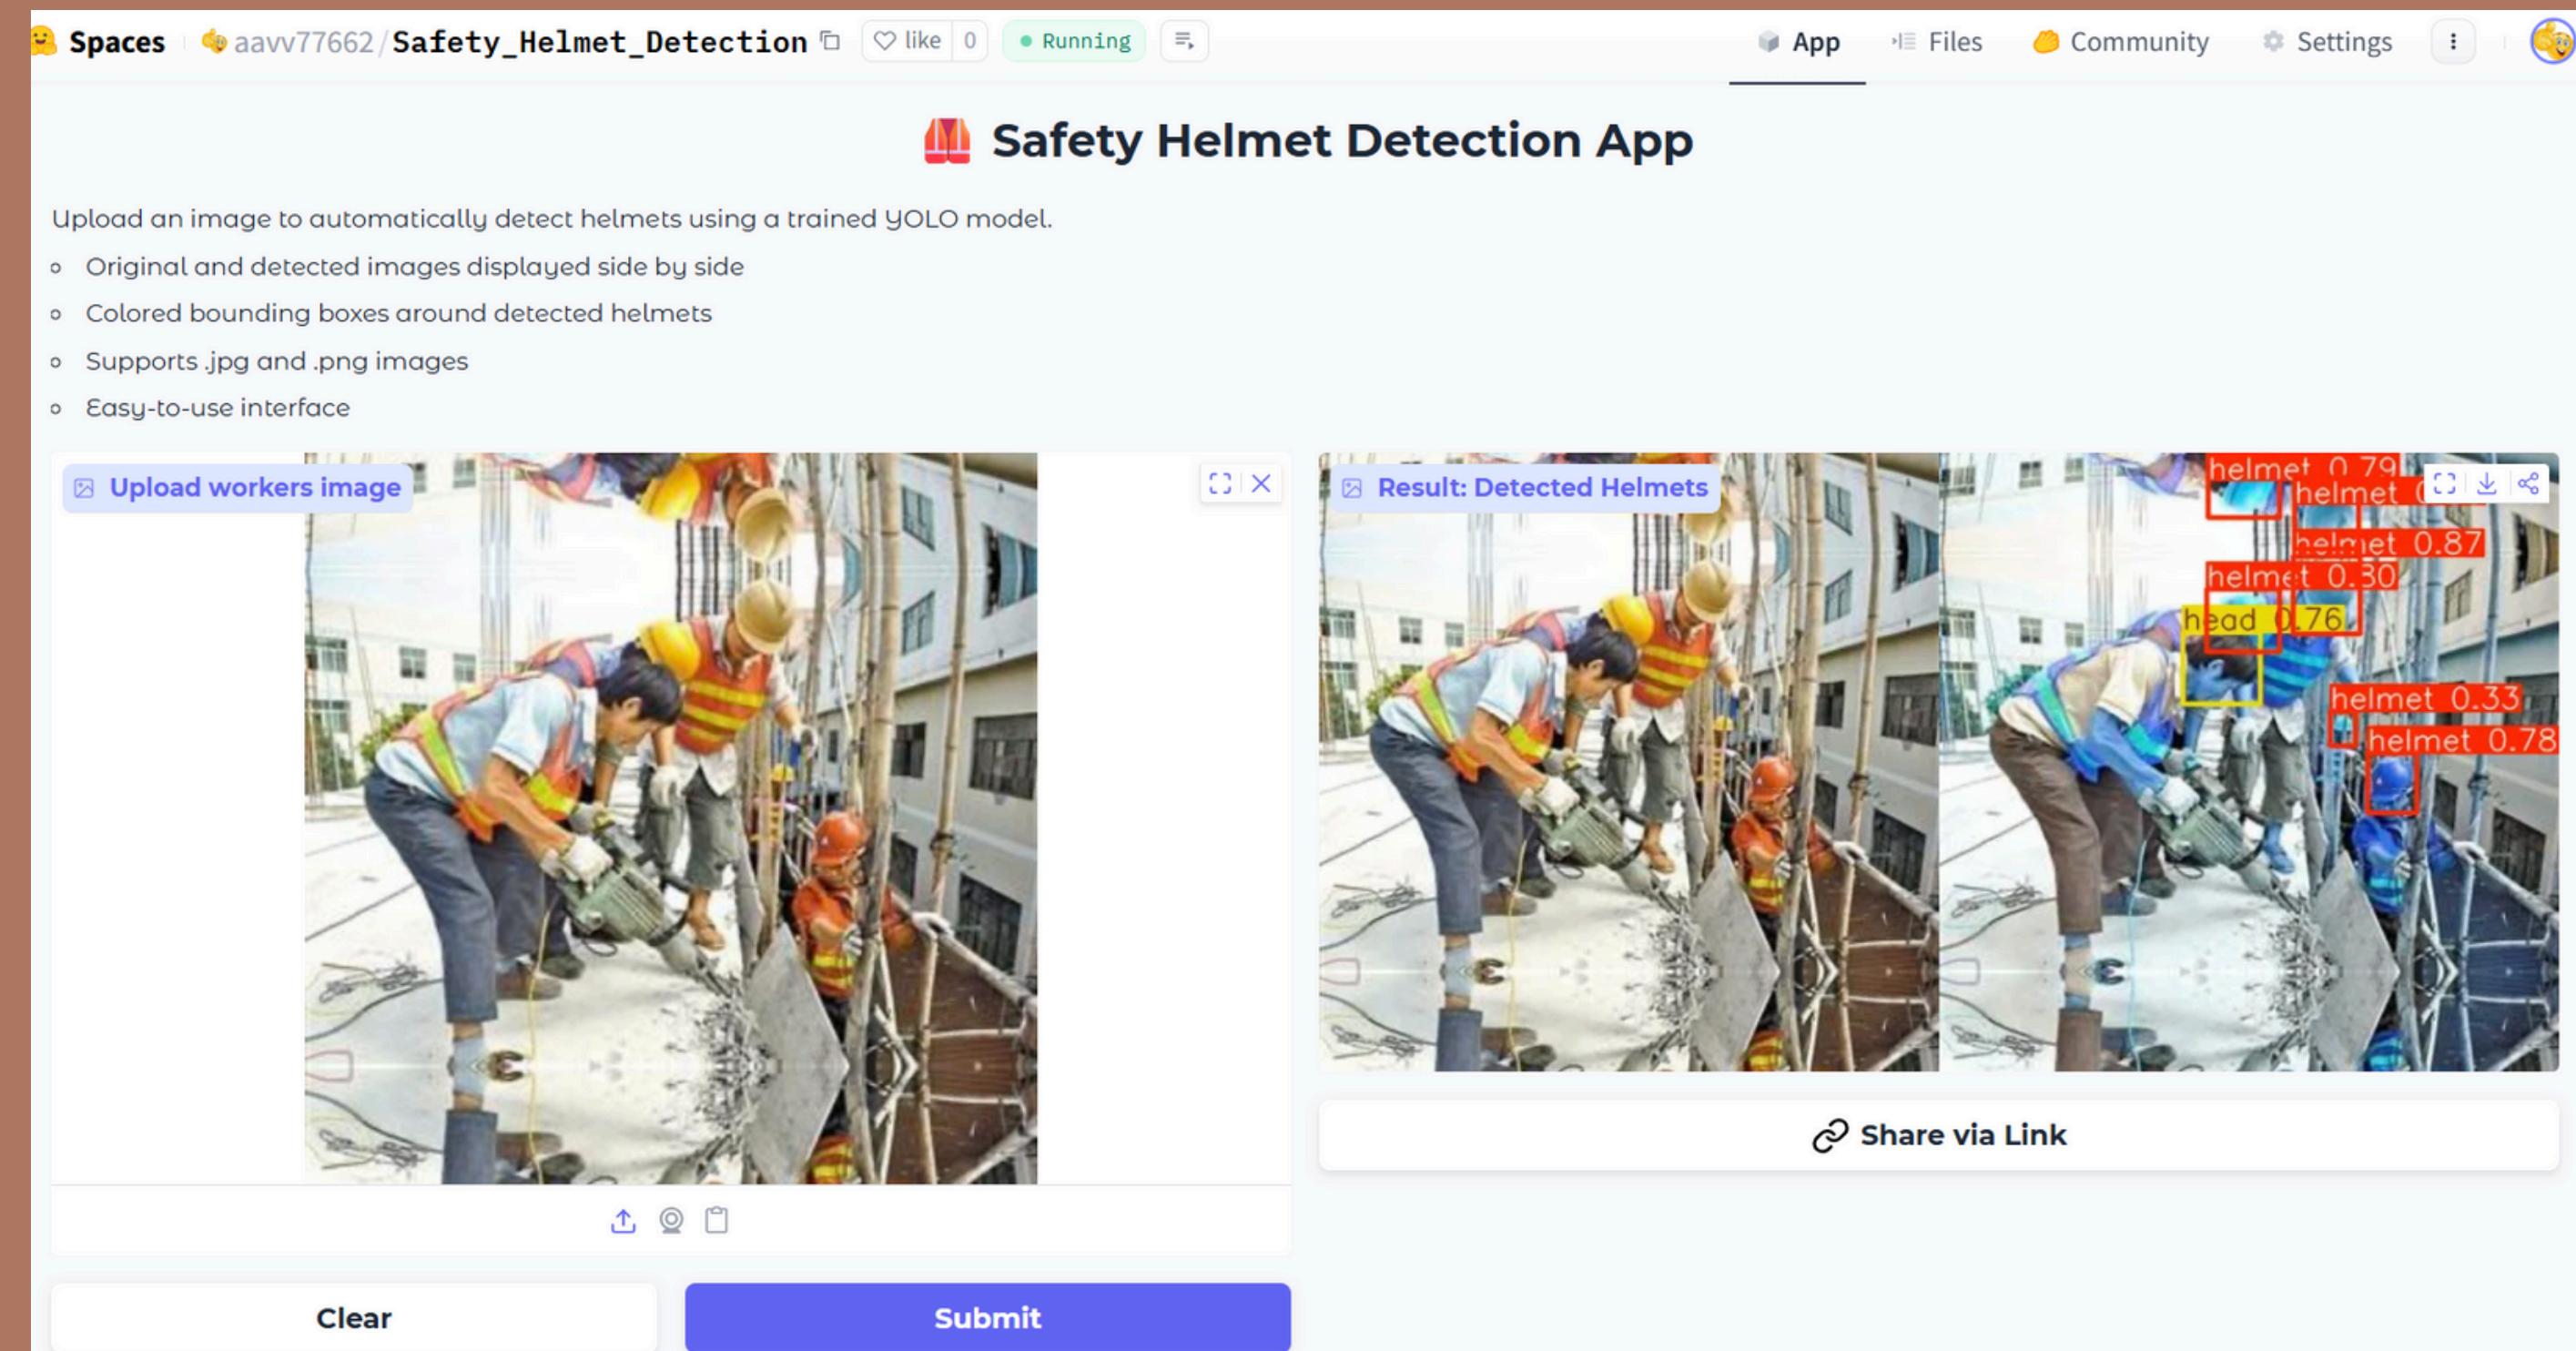


- Model uploaded to Hugging Face
- Easy access & sharing

Ready for:

- Web apps [https://huggingface.co/spaces/aavv77662/Safety\\_Helmet\\_Detection](https://huggingface.co/spaces/aavv77662/Safety_Helmet_Detection)
- APIs
- Cloud deployment

# Deployment



# Insight

This project demonstrates that lightweight deep learning models can effectively improve workplace safety when combined with proper data and system design.

## Limitations & Future Work

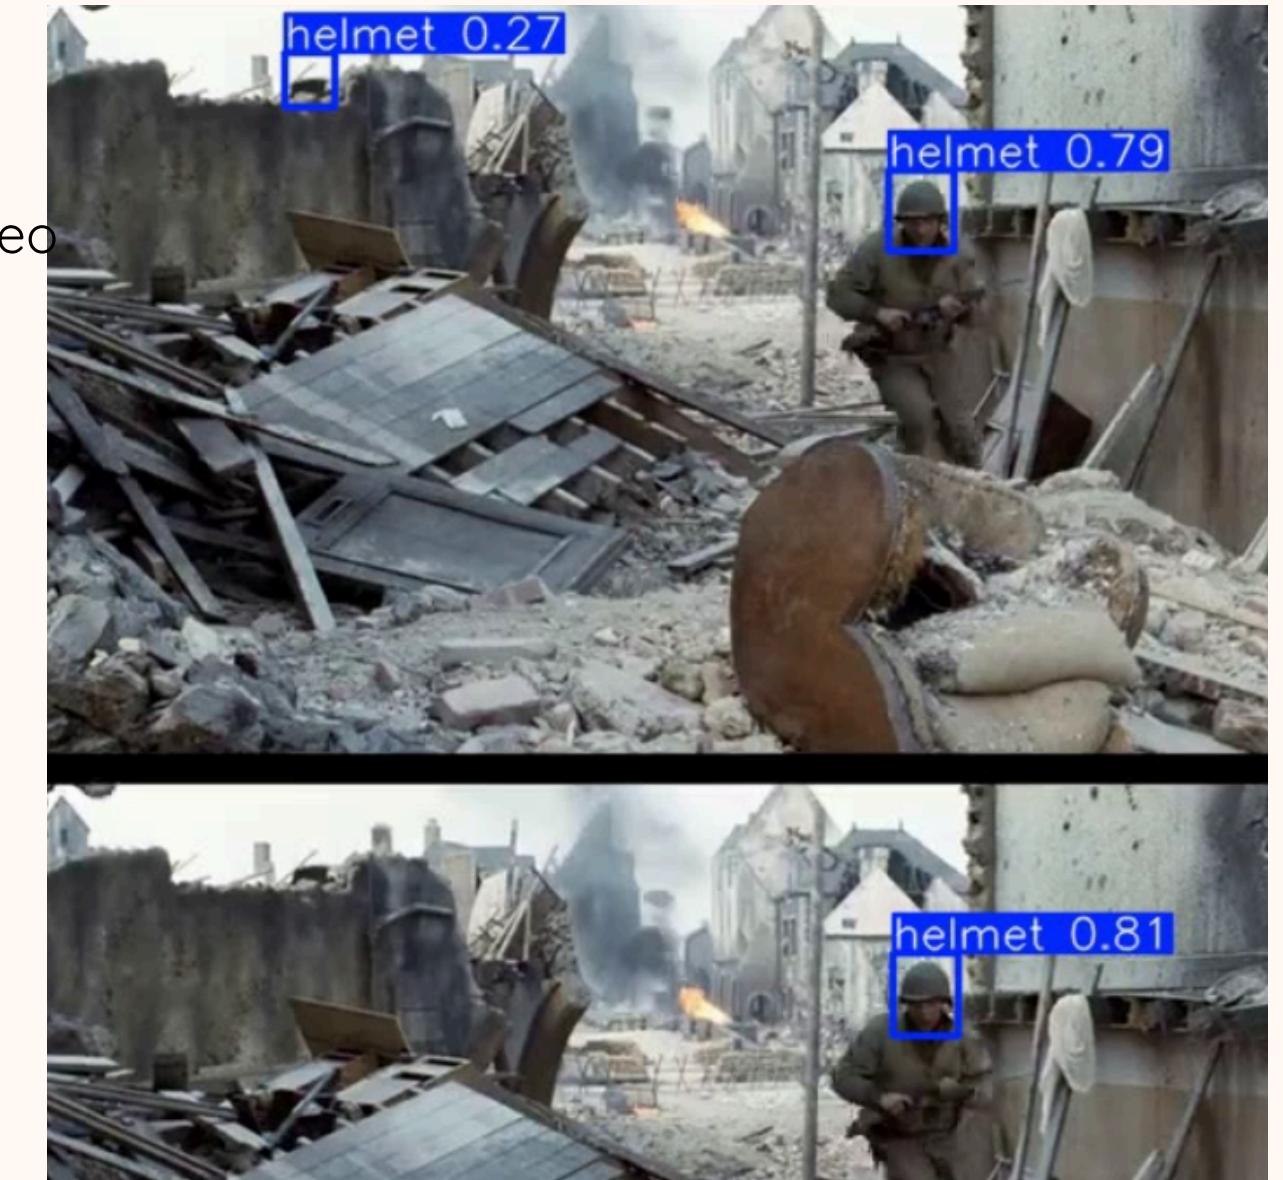
### Limitations

- Dataset imbalance
- Occlusion

### Future

- More data
- YOLOv8m
- Video + alerts

Output after test video  
to detection



# Thank You!

The project demonstrates an effective real-time helmet detection system using YOLOv8.